

APPS@UCU

Rust #4: Smart pointers, Concurrency and Async

Sultanov Andriy



Contents

- 1 Smart pointers
- 2 Concurrency
- 3 Looking back
- 4 And ahead...

Smart pointers

Smart pointers

General

Rust standard library has a few useful smart pointers, which modify the possibilities of operating with the value they hold:

- `Box<T>`
- `Rc<T>`
- `Cell<T>`
- `RefCell<T>`

Smart pointers

Deref trait

Implementing **Deref** allows Rust to treat smart pointers as references to the values they hold.

You can dereference types implementing **Deref** explicitly using the `*` operator, or you can omit the dereferencing and let Rust implicitly coerce your type to the type you need.

Thus, when calling a method on a type that implements **Deref** and does not implement the method, Rust will recursively dereference this type until it finds the appropriate method.

Smart pointers

Deref trait

So, let's say we want to have this kind of OOP behavior in Rust:

```
class Foo {  
    void m() { ... }  
}
```

```
class Bar extends Foo {}
```

```
public static void main(String[] args) {  
    Bar b = new Bar();  
    b.m();  
}
```

Smart pointers

Deref trait

We can model it like this:

```
use std::ops::Deref;
```

```
struct Foo {}
```

```
impl Foo {  
    fn m(&self) {  
        //..  
    }  
}
```

```
struct Bar {  
    f: Foo,  
}
```

```
impl Deref for Bar {  
    type Target = Foo;  
    fn deref(&self) -> &Foo {  
        &self.f  
    }  
}
```

Smart pointers

Deref trait

And then easily use it:

```
fn main() {  
    let b = Bar { f: Foo {} };  
    b.m();  
}
```

This allows you to easily operate with Rust's smart pointers.

Smart pointers

Deref trait

Rust will coerce types according to these rules:

- From `&T` to `&U` when `T: Deref<Target=U>`
- From `&mut T` to `&mut U` when `T: DerefMut<Target=U>`
- From `&mut T` to `&U` when `T: Deref<Target=U>`

Boxes are Rust's way to store data on the heap.

They are useful when you want to handle values whose sizes can not be known at compile time as if their size was known. This works since the size of the pointer to the heap that the Box keeps on the stack is known.

An example of moving a type that is typically stored on the stack onto the heap:

```
let val: u8 = 42;  
let boxed: Box<u8> = Box::new(val);
```

```
// We can easily dereference the box  
println!("{}", boxed); // prints 42
```

The value is going to be deallocated once it goes out scope, calling Box's **drop** implementation.

We could have just as well explicitly dereferenced the Box:

```
let val: u8 = 42;  
let boxed: Box<u8> = Box::new(val);
```

```
// We can easily dereference the box  
println!("{}", boxed); // prints 42
```

```
// And explicitly like this:  
println!("{}", *boxed);
```

Box<T>

If you want to create an unsized data structure, for example a recursive cons list, you can just use Box:

```
enum List<T> {  
    Cons(T, Box<List<T>>),  
    Nil,  
}
```

```
// Creating a new list
```

```
let list: List<i32> = List::Cons(1, Box::new(List::Cons(2, Box::new(List::Nil))));  
println!("{:?}", list); // prints Cons(1, Cons(2, Nil))
```

What will end up being stored is an **i32** and a pointer **usize** to another List on the heap.

While Rust's ownership system is pretty strict, there are some ways to have multiple ownership. `Rc<T>` is a reference counter that keeps track of the number of pointers to a certain value, and drops it once nobody uses it anymore.

Rc<T> does not allow mutability, and basically changes the ownership system so that the control over the lifetime of the value is done during runtime, not compile time.

Rc<T>

```
let text = "Rc examples".to_string();
{
    // Reference count is 1
    let rc_a: Rc<String> = Rc::new(text);
    {
        // Reference count is 2
        let rc_b: Rc<String> = Rc::clone(&rc_a);

        rc_a.len();
        println!("{}", rc_b);

        // rc_b is dropped, reference count is 1
    }
    // rc_a is dropped, reference count is 0, the value is dropped too
}
// Error, rc_examples was dropped!
println!("rc_examples: {}", rc_examples);
```


Rust also provides several other standard smart pointers, like `Cell<T>` and `RefCell<T>`, which allow mutating the immutable value in various cases.

We are not going to cover them in these lectures, but it's important to understand how all of these smart pointers work under the hood, and the way they allow us to operate with the values they hold - through `Deref`.

Concurrency

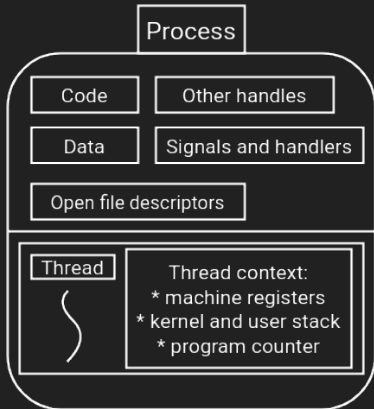
Concurrency in general

We've talked about how computers work in general before, but we've always implicitly assumed that they only run one thing at a time - that is, we've talked about a single-core CPU.

How are user programs represented for that CPU? Through the abstraction of a process - an instance of a program in execution. It might be better to think of the process as the collection of data structures that fully describes the state of the execution of a certain program.

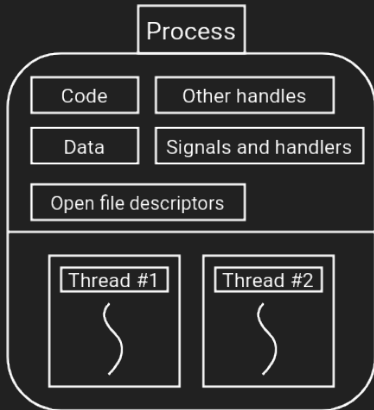
Concurrency in general

Thus, a process can roughly be represented as this, a collection of data structures that describe it, and a single thread of execution that is actually scheduled for execution and handles the state of the program.



Concurrency in general

One process can contain multiple threads of execution though, which are going to have the same shared memory and code through the process, but are going to be running different code simultaneously.



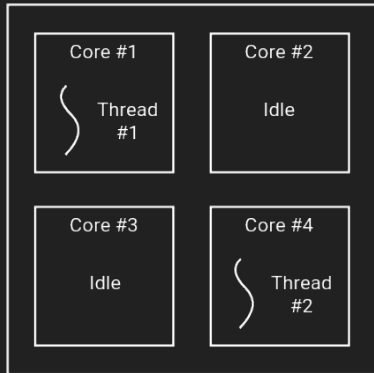
Concurrency in general

The task of scheduling and prioritizing these threads falls on the operating system, which allows creating these structures in the first place, and maintains a list of them, periodically switching the CPUs to run this or that thread.

We are not going to talk about multithreading from the OS perspective, but it's important to understand what happens with our CPU once we have multiple cores and are able to execute several threads at a time.

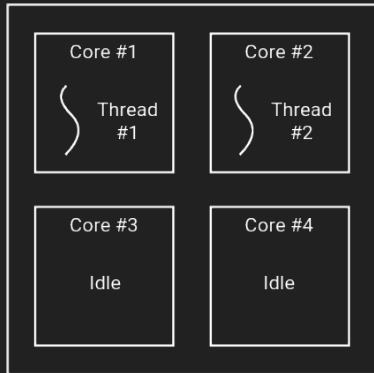
Concurrency in general

The operating system and the CPU handle all of the minor details of the process of scheduling and actually executing the threads on several cores, we are just going to assume that all of this works nicely (it almost always does)!



Concurrency in general

The operating system and the CPU handle all of the minor details of the process of scheduling and actually executing the threads on several cores, we are just going to assume that all of this works nicely (it almost always does)!



Synchronization

We've already talked about how a safe system would look like, where nobody will violate the memory safety of others:

- There can be many readers with no writers at the same time
- There can be only one writer with no readers at the same time
- Values can be used only as long as they still exist

We've seen how Rust ensures all of these rules are satisfied in a single-thread scenario, but what if there are several threads running at the same time? How do we make sure they don't write into each other's memory?

Concurrency

TODO: Talk about Arc, Mutex, mspc, and the way Rust's system prevents data races at compile time.

Looking back

Looking back

Over the course of four weeks we've covered some general topics, which I hope you will be able to apply everywhere, not only in Rust:

- Machine memory model
- Safe mutability rules
- Safe ownership and lifetime rules
- Multithreading

Looking back

And we've also looked at quite a lot of Rust-specific stuff:

- Rust's ownership/lifetime/mutability system
- Cargo ecosystem
- Algebraic types and error handling
- Declarative macros
- Traits
- Iterators and closures
- Multithreading primitives

And ahead...

Looking ahead

There is quite a lot of stuff in Rust that we haven't covered. If you want to dive deeper yourself, here is a short list of some of the most important things you can look into:

- Async Rust
- Procedural macros
- Testing
- FFI and bindings

And many more...

But, most importantly, practice makes perfect!

I hope that this short tour of Rust has shown you the possibilities of modern languages and ecosystems, and taught you a few important things in general!

Thank you!